

ECE 358 Assignment

Jason Shao, Lihao Luo, Minghao Lu

June 5, 2016

1. I am answering this question under the assumption that each edge corresponds to a unique router. I also assume that a routing table does not include an entry for itself as in the textbook, the example for hierarchical routing does not include itself as an entry (only adjacent hops).

Consider a network where the nodes (subnetworks) can form a Hamiltonian cycle (a cycle that includes all nodes in the graph). We can then take that cycle and create a routing table with exactly one entry that points to the next router (edge) in the cycle as the next hop for all destination addresses. Because we can create this cycle of all subnetworks, it is guaranteed that a packet forwarded in this cycle will be guaranteed to arrive at its intended subnetwork. For every other edge in the graph but not in the cycle, assign a single entry to any router that is directly in the cycle. This way, that router will forward the packet to a router in the Hamiltonian cycle which is then guaranteed to arrive at its destination. Here, we see that the lower bound is 1 entry if there exists a Hamiltonian cycle.

If it is the case that there is no Hamiltonian cycle, the lower bound of entries is the minimum of the maximum number incident edges for every spanning tree in G . Consider if you have a spanning tree of G and we used hierarchical routing such that all subnetworks are prefixed in a manner sensitive to the topology of the tree as we've seen from class and the slides. That means for any router, we only need to worry about the immediately incident edges as all other edges are hierarchally beneath those incident edges. Since the input graph is an arbitrary graph (not a tree), there may be multiple spanning trees that have different maximum number of incident edges for any edge. We choose the smallest value as the lower bound for the maximum table entry size. In the worst case scenario, it is possible that all other edges in the graph are incident to one single edge. That means the size of that router would be $n - 1$ since it must have an entry for every other router.

2. Yes, there exists a polynomial time algorithm for the decision problem. I will prove it by construction:
 - Look at every pair of prefixes
 - Check if any pair can be merged

- If so, remove the pair of prefixes and insert the aggregated prefix. Then re-start at step 1 by checking all pairs of prefixes again
- If not, check the final size of the list of prefixes, and compare it to k .

Checking if two prefixes can be merged can be done in $O(1)$ time as the IP address is bounded. Merging the two prefixes can also be done in $O(1)$ time by removing two existing 32-bit prefixes and inserting a new 32-bit prefix. There are $O(n^2)$ pairs to check where n is the initial number of prefixes. After we do a merge, we reduce the size of the set of prefixes and check all pairs again. Therefore, we can do at most $O(n)$ merges since it reduces the size of the set by 1 each time. In all, the algorithm has a worst case run-time of $O(n)O(n^2) = O(n^3)$. This is polynomial-time in the input size since there are n prefixes in the prefix list.

Correctness is guaranteed because the final prefix set is as minimal as it gets. This is because merging prefixes is a commutative operation (how you merge doesn't matter, the final result will be same). Since our algorithm keeps checking all combinations, when there are no more mergeable prefixes, there is no other way of merging to achieve a smaller set.

Assuming prefixes are stored as an array of bits, here is the pseudo-code:

```
def mergeable(prefix1, prefix2):
    if prefix1.prefix_length != prefix2.prefix_length:
        return False

    // Check that every bit is same except last bit
    for i <- 0 to prefix1.prefix_length - 2:
        if prefix1[i] != prefix2[i]:
            return False
    return prefix1[prefix_length - 1] != prefix2[prefix_length - 1]

def merge(prefix1, prefix2):
    // Decrement prefix length by 1
    output = prefix1
    output[prefix1.prefix_length - 1] = 0
    output.prefix_length = prefix1.prefix_length - 1

def decision(P, k):
    P' = P
    for i <- 0 to P'.size() - 1:
        for j <- i + 1 to P'.size() - 1:
            if mergeable(P'[i], P'[j]):
                P'.add(merge(P'[i], P'[j]))
```

```

P'.remove(P'[i])
P'.remove(P'[j])
i <- 0
j <- 0
return P'.size() <= k

```

3. In the site depicted on slide 20 of lecture 1-4, all traffic destined for 12.46.129.5 may arrive at interface D . This is because S is a multi-homed site with both provider-aggregatable (PA) and provider-independent (PI) IPv4 addresses. If the site uses the PA address, it advertises 12.46.129.0/25 at points C and D to $P1$ and $P2$ respectively. While $P1$ can aggregate this prefix into 12/8, $P2$ cannot since the prefix is not numerically adjacent to its own. For destinations matching 12.46.129.0/25, the prefix is more specific when going through $P2$ than $P1$. Therefore, traffic with destination 12.46.129.5 tends to go through $P2$ due to the longest-matching-prefix algorithm used in internet routing.

To address this, the site S can use the PI prefix instead of PA. Then, 198.134.135.0/24 is advertised to $P1$ and $P2$. Since neither of the ISPs are able to aggregate this prefix, the full 198.134.135.0/24 is advertised to both points A and B . Therefore, hosts would send datagrams to 198.134.135.0/24 and S can be reached naturally through whichever ISP is closer to the sender.

4. Yes, I concur with Alice.

Let T be the initial minimum spanning tree of G . Let u be the node that leaves G , and let v be the node such that uv is the only edge incident to u in T . Let T' be T without u and uv . I want to show that T' is the minimum spanning tree of the new graph.

Proof By Contradiction

Assume for the sake of contradiction that T' is not the most minimum spanning tree of the new graph. That means there must exist some tree R that spans the new graph, and the total weight of its edges is less than that of T' .

Now consider R with uv added to the tree. We can show that R with uv is a spanning tree of G , because R is a tree that spans every node in G except u . However, we know v is in R so uv can be added to R to span all nodes in G while maintaining the tree structure.

Moreover, observe that R with uv has total weight less than that of T because:

$$\begin{aligned}
& \text{weight}(R) < \text{weight}(T') \\
& \text{weight}(R) + \text{weight}(uv) < \text{weight}(T') + \text{weight}(uv) \\
& \text{weight}(R) + \text{weight}(uv) < \text{weight}(T)
\end{aligned}$$

Now we have shown that R with uv is a spanning tree of G and its weight is less than that of the T , the minimum spanning tree. This is clearly a contradiction, so we can declare that there does not exist another spanning tree that has total weight less than T' . Therefore, we do not need to recompute the MST as when a node leaves the network and it only has one edge incident on it, the new MST is merely the original tree with the node and its incident edge removed.

5. According to the book, for a request there are two "receiver" MAC address in the frame: DST address in the header that is 48 bits of 1 or Target address in the message that is just 0. I'm going to assume you want the latter.

part a)

(i) request (ii) 1.2.3.4 (iii) 1.2.3.4's MAC address (iv) 1.2.3.10 (v) 0
 (i) response (ii) 1.2.3.10 (iii) 1.2.3.10's MAC address (iv) 1.2.3.4 (v) 1.2.3.4's MAC address

part b)

(i) request (ii) 1.2.3.4 (iii) 1.2.3.4's MAC address (iv) 1.2.1.1 (v) 0
 (i) response (ii) 1.2.1.1 (iii) 1.2.1.1's MAC address (iv) 1.2.3.4 (v) 1.2.3.4's MAC address
 (i) request (ii) 10.11.12.1 (iii) 10.11.12.1's MAC address (iv) 10.11.12.25 (v) 0
 (i) response (ii) 10.11.12.25 (iii) 10.11.12.25's MAC address (iv) 10.11.12.1 (v) 10.11.12.1's MAC address
 (i) request (ii) 15.16.17.25 (iii) 15.16.17.25's MAC address (iv) 15.16.17.18 (v) 0
 (i) response (ii) 15.16.17.18 (iii) 15.16.17.18's MAC address (iv) 15.16.17.25 (v) 15.16.17.25's MAC address