ECE 358 Assignment 4

Jason Shao, Lihao Luo, Minghao Lu

June 23, 2016

1. (a) (A) 1.2.3.0/30
   (B) 1.2.3.1/30
   (C) 1.2.3.2/30
   (D) 1.2.3.3/30
   (E) 1.2.3.4/30
   (F) 1.2.3.5/30
   (G) 1.2.3.6/30
   (H) 1.2.3.7/30
   (I) 10.0.0.0/16
   (J) 10.0.0.1/16
   (K) 1.2.3.8/30
   (L) 1.2.3.9/30
   (M) 1.2.3.10/30
   (N) 1.2.3.11/30
   (O) 10.0.0.2/16
   (P) 10.0.0.3/16

   (b) (i)

Destination	Next Hop	Interface
1.2.3.4/30	myself	F
1.2.3.0/30	myself	C
0.0.0.0/0	1.2.3.1	C

   (ii)

Destination	Next Hop	Interface
1.2.3.8/30	myself	L
1.2.3.0/30	myself	D
0.0.0.0/0	1.2.3.1	D

2. During the first hop, the MTU is 1000 bytes, but the initial packet is $20 + 1800 = 1820$ bytes. Therefore, after fragmentation, f_1 will have 20 bytes of header and 976 bytes of payload since the offset has to be a multiple of 8 while maximizing the total packet size to less than 1000 bytes. Similarly, f_2 will have a header of 20 bytes and payload of the remaining 1800 - 976 = 824 bytes (offset of 122).

Afterwards, f_1 undergoes fragmentation again with MTU of 500 bytes. The first part, $f_1.1$ will have 20 bytes for the header and 480 bytes for payload. The second part, $f_1.2$ will have 20 bytes for the header and 480 bytes for payload (offset of 60). The third part, $f_1.3$ will have 20 bytes for the header and 976 - 480 - 480 = 16 bytes for the payload (offset of 120).

In conclusion, the final fragments received at the destination in order of offset is:

- First fragment: ID = abcd, More fragments = 1, Fragment offset = 0, Total length = 500 bytes (480 bytes of payload)
- Second fragment: ID = abcd, More fragments = 1, Fragment offset = 60, Total length = 500 bytes (480 bytes of payload)
- Third fragment: ID = abcd, More fragments = 1, Fragment offset = 120, Total length = 36 bytes (16 bytes of payload)
- Fourth fragment: ID = abcd, More fragments = 0, Fragment offset = 122, Total length = 844 bytes (824 bytes of payload)

3. (a) The header checksum is not necessarily the same, because the TTL field is decremented at each hop, so the header checksum is recomputed at each hop (hence, a different value than the initial checksum).

   (b) I do not concur, even if an odd number of bits are flipped, it is not guaranteed to detect an error. A counter-example is 01 FF (checksum FE) because if you flip all nine 1-bits into 00, the checksum is still FE.

   (c) Yes, the UDP checksum at the destination should match that of the source because UDP checksums are end-to-end and is not modified in transit (except when it passes through NAT).

   (d) No, the converse is "if the MTU is supported, you will always get a response" or "if you don't get a response, the MTU is not supported". This is not necessarily true, because there are other reasons for no response other than just a non-supported MTU (such as network congestion or TTL exceeded).

4. Like the assignment suggested, we adopt two premises. (i) at the point in time the slide considers, for every $a \in N'$, $D(a) = d(a)$. (ii) The path $u \rightsquigarrow y$ in the picture, $u \rightsquigarrow x \rightarrow y$, is a cheapest path from u to y. Let's call this path p_1.

First we prove $D(y) \leq cost(p_1)$. By definition $cost(p_1) = d(x) + c(x, y)$. Since $x \in N'$, by premise (i), $cost(p_1) = D(x) + c(x, y)$. When x was added to N', since y is adjacent to x, the algorithm performs $D(y) = min\{D(y), D(x) + c(x, y)\}$, and since the only operations performed on $D(y)$ is to assign it a min of its old value and another value, $D(y)$ never increases. Thus $D(y) \leq D(x) + c(x, y)$, which combined with $cost(p_1) = D(x) + c(x, y)$, implies $D(y) \leq cost(p_1)$.

Suppose $D(y) \neq d(y)$. Since $d(y)$ is the cheapest cost from u to y, and $D(y)$ is the cost of a path from u to y, $D(y) \geq d(y)$. Since $D(y) \neq d(y)$, $D(y) > d(y)$. Since

$d(y) < D(y)$ and $D(y) \leq cost(p_1)$, there must be a path from u to y that is cheaper than p_1. But premise (ii) says p_1 is a cheapest path from u to y, contradiction. Therefor $D(y) = d(y)$.

5. First we observe $\min_{v \in neigh(x)}\{c(x,v) + d_v(y)\}$ is at least an upper bound on $d_x(y)$. Let v_1 be a neighbour of x that achieves the minimum in $\min_{v \in neigh(x)}\{c(x,v) + d_v(y)\}$, i.e $c(x,v_1) + d_{v_1}(y) = \min_{v \in neigh(x)}\{c(x,v) + d_v(y)\}$. Let $v_1 \rightsquigarrow y$ be a minimum path from v_1 to y, i.e. $cost(v_1 \rightsquigarrow y) = d_{v_1}(y)$. Observe $x \rightarrow v_1 \rightsquigarrow y$ is a path from x to y. Moreover $cost(x \rightarrow v_1 \rightsquigarrow y) = c(x,v_1) + cost(v_1 \rightsquigarrow y) = c(x,v_1) + d_{v_1}(y) = \min_{v \in neigh(x)}\{c(x,v) + d_v(y)\})$. Since $d_x(y)$ is the cost of the cheapest path from x to y, $d_x(y) \leq cost(x, v_1 \rightsquigarrow y) = \min_{v \in neigh(x)}\{c(x,v) + d_v(y)\}$.

Suppose $d_x(y) = \min_{v \in neigh(x)}\{c(x,v) + d_v(y)\}$ is not true, since we proved $d_x(y) \leq \min_{v \in neigh(x)}\{c(x,v) + d_v(y)\}$, it must be the case $d_x(y) < \min_{v \in neigh(x)}\{c(x,v) + d_v(y)\}$. Let $p = x, v_1, v_2, ...v_k, y$ be a cheapest path from x to y, then $cost(p) = c(x,v_1) + cost(v_1, v_2, ..., v_k, y) < \min_{v \in neigh(x)}\{c(x,v) + d_v(y)\}$. In particular, $c(x,v_1) + cost(v_1, ..., v_k, y) < c(x,v_1) + d_{v_1}(y)$, which in turn implies $cost(v_1, ..., v_k, y) < d_{v_1}(y)$. Since $v_1, ..., v_k, y$ is a path from v_1 to y and $d_{v_1}(y)$ is the cost of the cheapest path from v_1 to y, this is a contradiction. Therefor our assumption, $d_x(y) = \min_{v \in neigh(x)}\{c(x,v) + d_v(y)\}$ is not true, must be false.

6. The number of nodes in G.

   Let $v_1, ..., v_k$ be the cheapest path between $v_1, ..., v_k$. First we observe $\forall i \leq k, v_1, v_2, ..., v_i$ is the cheapest path between v_1 and v_i. Suppose its not, then there must be path $v_1, w_1, w_2, ..., v_i$ that's cheaper, but then $v_1, w_1, w_2, ..., v_i, v_{i+1}, v_{i+2}, ..., v_k$ would be cheaper than $v_1, v_2, ..., v_k$. Since we defined $v_1, v_2, ..., v_k$ as the cheapest path between v_1 and v_k, this is a contradiction. Thus $\forall i \leq k : v_1, v_2, ..., v_i$ must be the cheapest path between v_1 and v_i.

   (For the following section, by "know" the cheapest path I mean know the next hop and total cost)

   In an iteration, if v_i doesn't know its cheapest path from v_1 yet, since v_i updates its guess of the cheapest from v_1 path based only on its neighbours guess of cheapest path from v_1, v_i would and would only learn its true cheapest path from v_1 if in the last iteration v_{i-1} knows its true cheapest path from v_1. Since by the end of the first iteration, only v_1 will know its distance for v_1, it takes a total of i iterations for v_i to learn its cheapest path from v_1. Thus it takes k iterations for v_k to learn its cheapest path from v_1.

   It follows that the number of iterations for everyone to know their cheapest path from everyone else will be the number of nodes in the longest path between any two nodes where that path is also the cheapest path between the two nodes. Such a path in a weighted, connected, undirected graph G could contain all the nodes in G.

3

7. In figure 1 below, there are 5 nodes in the undirected graph: A, B, C, D, E. The solid lines represent edges and the associated number is the weight of that undirected edge. The dashed line shows that node's minimum weight path to node E. The direction of the dashed line shows the next hop its current min path.

   i) The first diagram is the initial graph.

  ii) The event that causes the Distance Vector protocol to run is the change of weight on edge AE from 1 to 100. Node A is the first to realize this change. It is shown in the second diagram in the figure below.

 iii) Diagram 3 - 7 shows one iteration the routing loop, and the chain propagation of shortest path entry updates continues.

Figure 1: Example of routing loop of nodes A, B, C, D

Graph 1

E — 100 — A
8
4
9
1
1
1
B
4
D
C
10

Graph 2

E — 100 — A
8
4
9
1
1
1
B
11
D
C
10

Graph 3

E — 100 — A
15
4
9
1
1
1
B
11
D
C
10

and so on...